Automatic Sampling and Analysis of YouTube Data

Processing and Cleaning User Comments

Julian Kohne
Johannes Breuer
M. Rohangis Mohseni

2021-02-24

Processing and Cleaning User Comments

Preprocessing

- Preprocessing refers to all steps that need to be taken to make the data suitable for the actual analysis

- For webscraping data, this is often more tedious and time-consuming than for survey data because:

    - the data are not designed with your analysis in mind
    - the data are typically less structured
    - the data are typically more complex
    - the data are typically more heterogenous
    - the data are typically larger

- In addition, with large amounts of data it is often necessary to work on servers or clusters instead of regular desktop or laptop computers

- Even then, restructuring or transforming data can take days, so mistakes hurt more

Preprocessing

- In *Data Science*, most time is typically spent on the preprocessing rather than the actual analysis


What data scientists spend the most time doing

- Building training sets: 3%
- Cleaning and organizing data: 60%
- Collecting data sets; 19%
- Mining data for patterns: 9%
- Refining algorithms: 4%
- Other: 5%

Source: https://www.forbes.com/sites/gilpress/2016/03/23/data-preparation-most-time-consuming-least-enjoyable-data-science-task-survey-says/#157890a96f63

Preprocessing

- Also, it is perceived as the least enjoyable part of the process

What's the least enjoyable part of data science?

- Building training sets: 10%
- Cleaning and organizing data: 57%
- Collecting data sets: 21%
- Mining data for patterns: 3%
- Refining algorithms: 4%
- Other: 5%

Source: https://www.forbes.com/sites/gilpress/2016/03/23/data-preparation-most-time-consuming-least-enjoyable-data-science-task-survey-says/#157890a96f63

Preprocessing *YouTube* comments

- The `tuber` package returns an `R` dataframe instead of a JSON

- We can select which data we need by using the API through `tuber`

- For single videos, the data are small enough to be processed on a regular desktop/laptop computer

- However, this doesn't mean that the data are already usable for all intents and purposes

- We still need to:

  - select
  - format
  - extract
  - link

the information that is relevant to us

Preprocessing *YouTube* Comments

Loading the unprocessed comments into `R`

```r
# loading raw data
comments <- readRDS("../../data/RawEmojiComments.rds")
```

Understanding Your Data (1)

The first step is always to understand your data, this is especially crucial for *found data* because it was not designed with your analysis in mind

```
# listing all columns
colnames(comments)
```

```
##  [1] "videoId"              "textDisplay"            "textOriginal"
##  [4] "authorDisplayName"    "authorProfileImageUrl"  "authorChannelUrl"
##  [7] "authorChannelId.value" "canRate"               "viewerRating"
## [10] "likeCount"            "publishedAt"            "updatedAt"
## [13] "id"                   "parentId"               "moderationStatus"
```

Luckily, the *YouTube* API is very well documented and provides brief explanations for all the variables you can extract from it

Understanding Your Data (2)

This information is valuable for understanding what type of comments the dataframe contains

```
table(is.na(comments$parentId))
```

```
##
## FALSE   TRUE
## 15178 22474
```

A quick look at the documentation reveals:

parentID: *The unique ID of the parent comment. This property is only set if the comment was submitted as a reply to another comment.*

Understanding Your Data (3)

...or for knowing how specific data types are formatted

```
head(comments$publishedAt)
```

```
## [1] "2021-02-18T20:55:19Z" "2021-02-18T19:04:48Z" "2021-02-18T16:44:34Z"
## [4] "2021-02-18T02:42:39Z" "2021-02-18T00:28:40Z" "2021-02-17T21:54:55Z"
```

```
class(comments$publishedAt)
```

```
## [1] "character"
```

A quick look at the documentation reveals:

publishedAt: *The date and time when the comment was originally published. The value is specified in ISO 8601 (YYYY-MM-DDThh:mm:ss.sZ) format.*

Understanding Your Data (4)

...or how similar variables are different from each other

```
comments$textOriginal[195]
```

```
## [1] "1:03 Everyone When We First Saw This Was Happening"
```

```
strwrap(comments$textDisplay[195],79)
```

```
## [1] "<a href=\"https://www.youtube.com/watch?v=r8pJt4dK_s4&amp;t=1m03s\">1
## [2] "Everyone When We First Saw This Was Happening"
```

textOriginal: *The original, raw text of the comment as it was initially posted or last updated. The original text is only returned if it is accessible to the authenticated user, which is only guaranteed if the user is the comment's author.*

textDisplay: *The comment's text. The text can be retrieved in either plain text or HTML. (The comments.list and commentThreads.list methods both support a textFormat parameter, which specifies the desired text format). Note that even the plain text may differ from the original comment text. For example, it may replace video links with video titles.*

Selecting What You (Don't) need

Now we can decide on what we need for our analyses

```
Selection <- subset(comments,select = -c(authorProfileImageUrl,
                                         authorChannelUrl,
                                         authorChannelId.value,
                                         videoId,
                                         canRate,
                                         viewerRating,
                                         moderationStatus))
colnames(Selection)
```

```
## [1] "textDisplay"       "textOriginal"     "authorDisplayName"
## [4] "likeCount"         "publishedAt"      "updatedAt"
## [7] "id"                "parentId"
```

Word of advice: Always keep an unaltered copy of your raw data and don't overwrite it. You never know what kinds of mistakes/oversights you might notice down the line and you don't want to have to recollect everything. Save your parsed data in a separate file (or in multiple steps if your preprocessing pipeline is complex).

Formatting your Data

By default, the data you get out of `tuber` is most likely not in the right format for your analyses

```
sapply(Selection, class)
```

```
##       textDisplay      textOriginal authorDisplayName        likeCount
##       "character"       "character"       "character"      "character"
##       publishedAt         updatedAt                id         parentId
##       "character"       "character"       "character"      "character"
```

```
# Summary statistics for like counts
summary(Selection$likeCount)
```

```
##     Length      Class       Mode
##      37652  character  character
```

```
# time difference between first comment and now
Sys.time() - Selection$publishedAt[1]
```

```
## Error in unclass(e1) - e2: nicht-numerisches Argument für binären Operator
```

Formatting `likeCount`

We want the `likeCount` to be a numeric variable and the timestamps to be datetime objects

```r
# Transforming likeCount to numeric
# (NB: this overwrites the original column)
Selection$likeCount <- as.numeric(Selection$likeCount)

# testing
summary(Selection$likeCount)
```

```
##     Min.  1st Qu.   Median     Mean  3rd Qu.      Max.
##    0.000    0.000    2.000    4.551    5.000 4477.000
```

We can now work with the number of likes as a numeric variable

Formatting your Timestamps (1)

Timestamps are extremely complex objects due to:

- Different calendars
- Different formattings
- Different origins
- Different time zones
- Historical anomalies
- Different resolutions
- Summer vs. Wintertime (different for each country and depending on hemisphere!)
- Leap years
- etc.

For these reasons, **never** try to code your own time stamp translations from scratch. Fortunately, `R` has several build in methods for dealing with this madness. The most basic one is the `as.POSIXct()` function, the most convenient one is the `anytime()` function from the `anytime` package (another powerful option for dealing with times and dates in `R` is the `lubridate package` from the `Tidyverse`).

Formatting Timestamps (2)

```
# transforming timestamps to datetime objects
Selection$publishedAt[1]
```

```
## [1] "2021-02-18T20:55:19Z"
```

```
testtime <- as.POSIXct(Selection$publishedAt[1],
                       format = "%Y-%m-%dT%H:%M:%OSZ",
                       tz = "UTC")
testtime
```

```
## [1] "2021-02-18 20:55:19 UTC"
```

```
# testing whether we can compute a difference
# with the datetime object
Sys.time() - testtime
```

```
## Time difference of 5.059355 days
```

This internal representation of time objects will be extremely important for plotting trends over time and calculating time differences

Formatting Timestamps (3)

A more convenient way to transform datetimes is the `anytime package`. It automatically tries to guess the format from the character string, so you don't have to. This is especially handy for vectors of datetimes in multiple different formats.

```r
# transforming datetimes using anytime()
library(anytime)
Selection$publishedAt <- anytime(Selection$publishedAt,
                                 asUTC = TRUE)
Selection$updatedAt <- anytime(Selection$updatedAt,
                               asUTC = TRUE)
sapply(list(Selection$publishedAt,Selection$updatedAt),class)
```

```
##        [,1]       [,2]
## [1,] "POSIXct" "POSIXct"
## [2,] "POSIXt"  "POSIXt"
```

Word of Advice: For datetime conversions, always do some sanity checks, especially if you are using methods that automatically detect the format. Pay special attention to the *timezone* in which your data are saved and compare it to the documentation of the standard.

Formatting Timestamps (4)

Be aware of how to interpret your timestamps. Note that the date was interpreted as UTC but converted to our local CET timezone which is 1 hour ahead of UTC. This comment was made at 21:55:19 in *our time,* we have no idea about the time at the location of the user.

```
Selection$publishedAt[1]
```

```
## [1] "2021-02-18 21:55:19 CET"
```

Extracting Information

After having formatted all our selected columns, we usually also want to create some new columns with information that is not directly available in the raw data. For example, consider these comments:

```
# Example comments with extractable information
Selection$textOriginal[235]
```

```
## [1] "Amazing movie <U+0001F929>"
```

```
strwrap(Selection$textOriginal[34073],79)
```

```
## [1] "hey people! YOU can help! sign this!"
## [2] "http://www.ipetitions.com/petition/help-get-the-emoji-movie-canceled"
```

There are two issues exemplified by these comments:

1) Comments contain emoji and hyperlinks that might distort our text analysis later

2) These are features that we'd like to have in a separate column for our analysis

Extracting Hyperlinks (1)

We will start with deleting hyperlinks from our text and saving them in an additional column. We will use the text mining package qdapRegex for this as it has predefined routines for handling large text vectors and regular expressions.

```r
# Note that we are using the original text so we don't have
#to deal with the HTML formatting of the links
library(qdapRegex)
Links <- rm_url(Selection$textOriginal, extract = TRUE)
LinkDel <- rm_url(Selection$textOriginal)
Links[34071:34073]
```

```
## [[1]]
## [1] NA
##
## [[2]]
## [1] NA
##
## [[3]]
## [1] "http://www.ipetitions.com/petition/help-get-the-emoji-movie-canceled"
```

Extracting Hyperlinks (2)

We get back a list where each element corresponds to one row in the Selection dataframe and contains a vector of links that were contained in the textOriginal column. At the same time, the link was removed from the Selection dataframe.

```
LinkDel[34073]
```

```
## [1] "hey people! YOU can help! sign this!"
```

```
Links[34073]
```

```
## [[1]]
## [1] "http://www.ipetitions.com/petition/help-get-the-emoji-movie-canceled"
```

Extracting Emoji (1)

The `qdapRegex` package has a lot of other different predefined functions for extracting or removing certain kinds of strings:

- `rm_citation()`
- `rm_date()`
- `rm_phone()`
- `rm_postal_code()`
- `rm_email()`
- `rm_dollar()`
- `rm_emoticon()`

Unfortunately, it does **not** contain a predefined method for emoji, so we will have to use the `emo` package for removing the emoji and come up with our own method for extracting them.

Extracting Emoji (2)

First we want to replace the emoji with a textual description, so that we can treat it just like any other token in text mining. This is no trivial task, as we have to go through each comment and replace each emoji with its respective textual description. Unfortunately, we did not find a working, easy to use, out of the box solution for this. But we can always make our own!

Essentially, we want to replace this:

```
emo::ji("monkey")
```

```
## <U+0001F412>
```

with this

```
"EMOJI_Monkey"
```

```
## [1] "EMOJI_Monkey"
```

Extracting Emoji (3)

First of all, we need a dataframe that contains the emoji as they are internally represented by R (this can be quite the hassle). Luckily, this is contained in the emo package.

```r
library(emo)
EmojiList <- jis
EmojiList[1:3,c(1,3,4)]
```

```
## # A tibble: 3 x 3
##   runes emoji       name
##   <chr> <chr>       <chr>
## 1 1F600 "\U0001f600" grinning face
## 2 1F601 "\U0001f601" beaming face with smiling eyes
## 3 1F602 "\U0001f602" face with tears of joy
```

Extracting Emoji (4)

Next, we need to paste the names of the emoji together while capitalizing the first letter of every word for better readability

```r
# Defining a function for capitalizing and pasting names together
simpleCap <- function(x) {

  # Splitting the string
  splitted <- strsplit(x, " ")[[1]]

  # Pasting it back together with capital letters
  paste(toupper(substring(splitted, 1,1)),
        substring(splitted, 2),
        sep = "",
        collapse = " ")
}
```

Extracting Emoji (5)

```r
# Applying the function to all the names
CamelCaseEmojis <- lapply(jis$name, simpleCap)
CollapsedEmojis <- lapply(CamelCaseEmojis,
                          function(x){gsub(" ",
                                           "",
                                           x,
                                           fixed = TRUE)})

EmojiList[,4] <- unlist(CollapsedEmojis)
EmojiList[1:3,c(1,3,4)]
```

```
## # A tibble: 3 x 3
##   runes emoji        name
##   <chr> <chr>        <chr>
## 1 1F600 "\U0001f600" GrinningFace
## 2 1F601 "\U0001f601" BeamingFaceWithSmilingEyes
## 3 1F602 "\U0001f602" FaceWithTearsOfJoy
```

Extracting Emoji (6)

After that, we need to order our dictionary from longest to shortest, so that we can prevent partial matching of shorter strings later

```
EmojiList <- EmojiList[rev(order(nchar(jis$emoji))),]
head(EmojiList[,c(1,3,4)],5)
```

```
## # A tibble: 5 x 3
##   runes                      emoji                                        name
##   <chr>                      <chr>                                        <chr
## 1 1F469 200D 2764 FE0F 200~  "\U0001f469<U+200D><U+2764><U+FE0F><U+200D>\U0
## 2 1F468 200D 2764 FE0F 200~  "\U0001f468<U+200D><U+2764><U+FE0F><U+200D>\U0
## 3 1F469 200D 2764 FE0F 200~  "\U0001f469<U+200D><U+2764><U+FE0F><U+200D>\U0
## 4 1F3F4 E0067 E0062 E0077 ~  "\U0001f3f4\U000e0067\U000e0062\U000e007~ Wale
## 5 1F3F4 E0067 E0062 E0073 ~  "\U0001f3f4\U000e0067\U000e0062\U000e007~ Scot
```

Note that what we are ordering by the `emoji` column, not the `text` or `runes` columns

Extracting Emoji (7)

Now we can `loop` through all of our emoji and replace them one after the other in each comment (*note*: this may take a while)

```r
# Assigning the column to a TextEmoRep variable
TextEmoRep <- LinkDel

# Looping through all Emojis for all comments in LinkDel
for (i in 1:dim(EmojiList)[1]) {

  TextEmoRep <- rm_default(TextEmoRep,
                pattern = EmojiList[i,3],
                replacement = paste0("EMOJI_",
                                EmojiList[i,4],
                                " "),
                fixed = TRUE,
                clean = FALSE,
                trim = FALSE)
}
```

Extracting Emoji (8)

As output, we get a large character vector with replaced emoji

```
Selection$textOriginal[235]
```

```
## [1] "Amazing movie <U+0001F929>"
```

```
TextEmoRep[235]
```

```
## [1] "Amazing movie EMOJI_Star-struck "
```

Extracting Emoji (9)

```r
ExtractEmoji <- function(x){

  SpacerInsert <- gsub(" ","[{[SpAC0R]}]", x)
  ExtractEmoji <- rm_between(SpacerInsert,
                             "EMOJI_","[{[SpAC0R]}]",
                             fixed = TRUE,
                             extract = TRUE,
                             clean = FALSE,
                             trim = FALSE,
                             include.markers = TRUE)

  UnlistEmoji <- unlist(ExtractEmoji)
  DeleteSpacer <- sapply(UnlistEmoji,
                    function(x){gsub("[{[SpAC0R]}]",
                                     " ",
                                     x,
                                     fixed = TRUE)})

  names(DeleteSpacer) <- NULL
  Emoji <- paste0(DeleteSpacer, collapse = "")
  return(Emoji)
}
```

Extracting Emoji (10)

We can apply the function to get one vector containing only the emoji as textual descriptions

```
Emoji <- sapply(TextEmoRep,ExtractEmoji)
names(Emoji) <- NULL
LinkDel[235]
```

```
## [1] "Amazing movie <U+0001F929>"
```

```
Emoji[235]
```

```
## [1] "EMOJI_Star-struck "
```

Removing Emoji

In addition, we remove the emoji from our `LinkDel` variable to have one *clean* column that we can use for text mining later. This column will not contain hyperlinks or emoji.

```
# We take the LinkDel column and also delete the emoji from it
library(emo)
LinkDel[235]
```

```
## [1] "Amazing movie <U+0001F929>"
```

```
TextEmoDel <- ji_replace_all(LinkDel,"")
TextEmoDel[235]
```

```
## [1] "Amazing movie "
```

Extracting Information

We now have different versions of our text column

1) The original one, with hyperlinks and emoji (`Selection$textOriginal`)

2) One with only plain text and without hyperlinks and emoji (`TextEmoDel`)

3) One with only hyperlinks (`Links`)

4) One with only emoji (`Emoji`)

We want to integrate all of them in our dataframe.

Linking everything back together

We can now combine our dataframe with the additional columns we created to have the perfect starting point for our analysis! However, because we sometimes have more than two links or two emoji per comment, we need to use the `I()` function so we can put them in the dataframe as `is`. Later, we will have to unlist these columns rowwise if we want to use them.

```r
df <- cbind.data.frame(Selection$authorDisplayName,
                       Selection$textOriginal,
                       TextEmoRep,
                       TextEmoDel,
                       Emoji = I(Emoji),
                       Selection$likeCount,
                       Links = I(Links),
                       Selection$publishedAt,
                       Selection$updatedAt,
                       Selection$parentId,
                       Selection$id,
                       stringsAsFactors = FALSE)
```

Linking everything back together

As a final step, we can give the columns appropriate names and save the dataframe for later use

```r
# setting column names
names(df) <- c("Author",
               "Text",
               "TextEmojiReplaced",
               "TextEmojiDeleted",
               "Emoji",
               "LikeCount",
               "URL",
               "Published",
               "Updated",
               "ParentId",
               "CommentID")

saveRDS(df, file = "../../data/ParsedEmojiComments.rds")
```

Exercise time 🏋️ 💪 🏃 🚴

Solutions